

# Deep Learning II: Optimization

CSci 5525: Machine Learning


Instructor: Nicholas Johnson

October 15, 2020

# Announcements

- HW2 is due tonight 11:59 PM CDT
- Exam 1 will be posted on Oct 20 (due 48 hours later)
- Course feedback form
- Project proposal grades/feedback posted

# The Picture: General




- Between input and hidden layers

$$w_{ji}^{(1)} \rightarrow \underline{a_j} = \sum_{i=0}^D w_{ji} x_i \quad \underline{z_j} = h(a_j)$$

- Between hidden and output layers

$$w_{kj}^{(2)} \rightarrow \underline{a_k} = \sum_{j=0}^M w_{kj} z_j \quad \underline{\hat{y}_k} = h(a_k)$$

# The Picture: General



$$\hat{y}_k(\mathbf{x}, \mathbf{w}) = \sigma \left( \sum_{j=0}^M w_{kj}^{(2)} h \left( \sum_{i=0}^D w_{ji}^{(1)} x_i \right) \right)$$

- Consider ERM problem with squared loss

$$\operatorname{argmin}_{\theta} \frac{c}{n} \sum_{i=1}^n \|\mathbf{y}_i - \hat{\mathbf{y}}_i\|_2^2$$


where  $\hat{\mathbf{y}}_i = f(\mathbf{x}_i, \theta)$  is the predicted output of the network and  $\mathbf{y}_i$  is the target vector and  $c > 0$  is a constant

- For the  $i$ th sample, the error is  $E_i = \frac{1}{2} \sum_{k=1}^K (y_k - \hat{y}_k)^2$  where  $\hat{y}_k$  is the  $k$ th output node value and  $y_k$  is the  $k$ th target value

# Training

- The process of sending a sample through the neural network to compute the prediction is called the forward propagation step
- To train the network, we use an algorithm called **backpropagation** (backprop) ←
- Backprop is a message-passing algorithm which attributes credit of the error to different nodes depending on their values  
→ • Nodes which have higher activation contribute more and are assigned higher credit for the error
- Backprop is an algorithm that computes the chain rule of calculus
  - Chain rule is used to compute derivatives of functions formed by composing other functions whose derivatives are known

# The Picture: General



$$\sigma(a_k) = h(a_k)$$

$$\hat{y}_k(\mathbf{x}, \mathbf{w}) = \sigma \left( \sum_{j=0}^M w_{kj}^{(2)} h \left( \sum_{i=0}^D w_{ji}^{(1)} x_i \right) \right)$$

# Backpropagation Derivation

→ The squared error on a single example is  $E = \frac{1}{2} \sum_k (y_k - \hat{y}_k)^2$  ←  
Compute gradients to propagate errors back through second layer:

$$\begin{aligned}\frac{\partial E}{\partial w_{k,j}} &= \underbrace{-(y_k - \hat{y}_k)}_{\text{error}} \underbrace{\frac{\partial \hat{y}_k}{\partial w_{k,j}}}_{\text{gradient of output w.r.t weight}} \\ &= -(y_k - \hat{y}_k) \frac{\partial h(a_k)}{\partial w_{k,j}} \\ &= \underbrace{-(y_k - \hat{y}_k)}_{\text{error}} \underbrace{h'(a_k)}_{\text{activation derivative}} \underbrace{\frac{\partial a_k}{\partial w_{k,j}}}_{\text{gradient of net input w.r.t weight}} \\ &= -(y_k - \hat{y}_k) h'(a_k) \frac{\partial}{\partial w_{k,j}} \left( \sum_j w_{k,j} z_j \right) \\ &= -(y_k - \hat{y}_k) h'(a_k) z_j \\ &\rightarrow \boxed{= -z_j \times \Delta_k}\end{aligned}$$

$\hat{y}_k = h(a_k)$

where  $\Delta_k = (y_k - \hat{y}_k) \times h'(a_k)$  and  $z_j$  is  $j$ th hidden unit



# Backprop: Output Layer Update Equations

- 
- $\frac{\partial E}{\partial w_{k,j}} = -z_j \times \Delta_k$
  - Use gradients in gradient descent:  $w_{k,j} \leftarrow w_{k,j} - \alpha \frac{\partial E}{\partial w_{k,j}}$
  - Backpropagate error to output layer:
- $$w_{k,j} \leftarrow w_{k,j} + \alpha \times z_j \times \Delta_k$$

where  $\Delta_k = (y_k - \hat{y}_k) \times h'(a_k)$  and  $\alpha$  is learning rate

# Backpropagation Derivation (cont.)

$$\begin{aligned}\frac{\partial E}{\partial w_{j,i}} &= - \sum_k (y_k - \hat{y}_k) \frac{\partial \hat{y}_k}{\partial w_{j,i}} = - \sum_k (y_k - \hat{y}_k) \frac{\partial h(a_k)}{\partial w_{j,i}} \\&= - \sum_k \underbrace{(y_k - \hat{y}_k) h'(a_k)}_{\Delta_k} \frac{\partial a_k}{\partial w_{j,i}} = - \sum_k \Delta_k \frac{\partial}{\partial w_{j,i}} \left( \sum_j w_{k,j} z_j \right) \\&= - \sum_k \Delta_k \underbrace{w_{k,j}}_{\text{hidden units}} \frac{\partial z_j}{\partial w_{j,i}} = - \sum_k \Delta_k w_{k,j} \frac{\partial h(a_j)}{\partial w_{j,i}} \\&= - \sum_k \Delta_k w_{k,j} h'(a_j) \frac{\partial a_j}{\partial w_{j,i}} \\&= - \sum_k \Delta_k w_{k,j} h'(a_j) \frac{\partial}{\partial w_{j,i}} \left( \sum_i w_{j,i} x_i \right) \\&= - \sum_k \Delta_k w_{k,j} h'(a_j) \underline{x_i} = \boxed{-x_i \times \Delta_j}\end{aligned}$$

where  $\Delta_j = h'(a_j) \sum_k w_{k,j} \Delta_k$

# Backprop: Hidden Layer Update Equations


- $\frac{\partial E}{\partial w_{j,i}} = -x_i \times \Delta_j$

- Use gradients in gradient descent:  $w_{j,i} \leftarrow w_{j,i} - \alpha \frac{\partial E}{\partial w_{j,i}}$
- Backpropagate error from the output layer to hidden layer:

$$w_{j,i} \leftarrow w_{j,i} + \alpha \times x_i \times \Delta_j$$

where  $\Delta_j = h'(a_j) \sum_k w_{k,j} \Delta_k$

# Challenges in Deep Learning Optimization

- 
- Local minima, weight space symmetry
  - Plateaus, saddle points, flat regions
  - Cliffs and exploding gradients
  - Long term dependence
  - Inexact gradients
  - Choice of step size
  - Scalability, first-order and second-order methods

# Cliffs and Exploding Gradients




Figure 8.3: The objective function for highly nonlinear deep neural networks or for recurrent neural networks often contains sharp nonlinearities in parameter space resulting from the multiplication of several parameters. These nonlinearities give rise to very high derivatives in some places. When the parameters get close to such a cliff region, a gradient descent update can catapult the parameters very far, possibly losing most of the optimization work that had been done. Figure adapted with permission from [Pascanu et al. \(2013\)](#).

# Exploding and Vanishing Gradients



- For an  $L$  layer neural network

$$\begin{aligned} \rightarrow f &= f_L \circ f_{L-1} \dots \circ f_1 \\ f(x) &= f_L(f_{L-1}(\dots f_1(x) \dots)) \end{aligned}$$

- Gradient (Jacobian) of  $f$  with respect input  $x$  is the product

$$f' = f'_L f'_{L-1} \dots f'_1 \quad \Bigg)$$

- Consequence: Exploding or vanishing gradients

# Optimization

- Mini-batch updates: Sample  $\underline{m}$  examples  $\{x_i\}$

$$\hat{g}_t \leftarrow \nabla_{\theta} \frac{1}{m} \sum_{i=1}^m \ell(f(\mathbf{x}_i; \theta_t), y_i)$$
$$\theta_{t+1} \leftarrow \theta_t - \eta_t \hat{g}_t$$

$\downarrow$   $\downarrow$   $\downarrow$   $1 \leq m \leq n$   $O\left(\frac{m}{\epsilon}\right)$

- • Momentum methods
- • Adaptive gradient descent
- • Combinations of such ideas

# Momentum Method




Figure 8.5: Momentum aims primarily to solve two problems: poor conditioning of the Hessian matrix and variance in the stochastic gradient. Here, we illustrate how momentum overcomes the first of these two problems. The contour lines depict a quadratic loss function with a poorly conditioned Hessian matrix. The red path cutting across the contours indicates the path followed by the momentum learning rule as it minimizes this function. At each step along the way, we draw an arrow indicating the step that gradient descent would take at that point. We can see that a poorly conditioned quadratic objective looks like a long, narrow valley or canyon with steep sides. Momentum correctly traverses the canyon lengthwise, while gradient steps waste time moving back and forth across the narrow axis of the canyon. Compare also figure 4.6, which shows the behavior of gradient descent without momentum.



# SGD with Momentum

- Velocity variable  $v$ 
  - Direction + ~~velocity~~ <sup>speed</sup> at which the parameters move through space
  - Dynamical system perspective
- SGD/Mini-batch update with momentum

$$\left( \begin{array}{l} \hat{g}_t \leftarrow \nabla_{\theta} \frac{1}{m} \sum_{i=1}^m \ell(f(\mathbf{x}_i; \theta_t), y_i) \\ \rightarrow v_{t+1} \leftarrow \alpha_t v_t - \eta_t \hat{g}_t \\ \theta_{t+1} \leftarrow \theta_t + v_{t+1} \end{array} \right. \quad \theta_{t+1} = \theta_t - \eta_t \hat{g}_t$$

- Parameter  $\alpha \in [0, 1)$  determines decay rate of gradient
- Dynamics of  $\alpha$ : start small, increase over iterations

# SGD with Nesterov Momentum

- Inspired by Nesterov's accelerated gradient descent
- Evaluate gradient at  $\theta_t + v_t$  instead of at  $\theta_t$
- SGD with Nesterov momentum

$$\begin{aligned} \tilde{\theta}_t &\leftarrow \theta_t + \alpha_t v_t \\ \hat{g}_t &\leftarrow \nabla_{\tilde{\theta}} \frac{1}{m} \sum_{i=1}^m \ell(f(\mathbf{x}_i; \tilde{\theta}_t), y_i) \\ v_{t+1} &\leftarrow \alpha_t v_t - \eta_t \hat{g}_t \\ \theta_{t+1} &\leftarrow \theta_t + v_{t+1} \end{aligned}$$

# Adaptive Learning Rate

- Adapt the learning rate over  $t$

$$\theta_{t+1} = \theta_t - \underline{\eta_t} \nabla \ell(\theta_t)$$

in: 1. LR



- Basic GD/SGD have decreasing learning rate:  $\eta_t = \frac{\eta_0}{\sqrt{\gamma t}}$ ,  
 $\eta_t = \frac{\eta_0}{\gamma t}$ , etc.
- Limitation: Exact same learning rate across all dimensions, slow convergence
- Goal: different/adaptive learning rates for different parameters

# Adaptive Learning Rate: Adagrad

- Gradient of the  $i$ th parameter  $\theta^i$  at iteration  $t$ :

$$\longrightarrow g_t^i = \nabla_{\theta^i} \ell(\theta_t)$$

- The adaptive learning rate used by Adagrad

$$\longrightarrow \eta_t^i = \frac{\eta_0}{\sqrt{\sum_{s=0}^t (g_s^i)^2}}$$

$$d\left(\frac{1}{\sqrt{t}}\right)$$

- SGD/minibatch with Adagrad

$$\longrightarrow \hat{g}_t \leftarrow \nabla_{\tilde{\theta}} \frac{1}{m} \sum_{i=1}^m L(f(\mathbf{x}_i; \tilde{\theta}_t), y_i)$$

$$\longrightarrow r_t \leftarrow r_{t-1} + \hat{g}_t \odot \hat{g}_t$$

small  
value

$$\Delta \theta \leftarrow - \frac{\eta_0}{\delta + \sqrt{r_t}} \odot \hat{g}_t$$

$$\longrightarrow \theta_{t+1} \leftarrow \theta_t + \Delta \theta$$

# Adagrad: Advantages

- 
- Parameters with large gradients  $\Rightarrow$  small learning rate and small gradients  $\Rightarrow$  large learning rate.
  - Suitable for deep learning
    - Scale of the gradient for different  $\theta^i$  often different by several orders of magnitude
    - Adjust step size based on scale
  - Empirical results on large-scale problems are good

# Adagrad: Disadvantages

- If initial gradients are large, subsequent training will be slow
- Learning rate monotonically decreases
  - May lead to early training termination
  - Maybe problematic for deep learning, e.g., exploding gradients
- Hyper-parameter  $\eta_0$  needs to be chosen

# Adaptive Learning Rate: RMSprop

- Need to recover (increase) step size over 'flat' regions
- Uses an exponentially decaying average for  $r$  (squared gradient)
- RMSProp algorithm

Handwritten annotations for the RMSProp algorithm equations:

- momentum with an arrow pointing to the  $\rho$  term in the  $r_t$  equation.
- adaptive with an arrow pointing to the  $\hat{g}_t \odot \hat{g}_t$  term in the  $r_t$  equation.
- $1e^{-5}$  with an arrow pointing to the  $\eta_0$  term in the  $\Delta\theta$  equation.

$$\begin{aligned}\hat{g}_t &\leftarrow \nabla_{\tilde{\theta}} \frac{1}{m} \sum_{i=1}^m \ell(f(\mathbf{x}_i; \tilde{\theta}_t), y_i) \\ r_t &\leftarrow \rho r_{t-1} + (1 - \rho) \hat{g}_t \odot \hat{g}_t \\ \Delta\theta &\leftarrow - \frac{\eta_0}{\sqrt{\delta + r_t}} \odot \hat{g}_t \\ \theta_{t+1} &\leftarrow \theta_t + \Delta\theta\end{aligned}$$

- Extension to Nesterov momentum: gradient at  $\tilde{\theta}_t = \theta_t + \alpha_t v_t$


# Adaptive Learning Rate: Adam

- Idea: RMSprop + first and second order momentum + bias correction
- Adam algorithm


$$\begin{aligned}\hat{g}_t &\leftarrow \nabla_{\tilde{\theta}} \frac{1}{m} \sum_{i=1}^m L(f(\mathbf{x}_i; \tilde{\theta}_t), y_i) \\ s_t &\leftarrow \rho_1 s_{t-1} + (1 - \rho_1) \hat{g}_t \\ r_t &\leftarrow \rho_2 r_{t-1} + (1 - \rho_2) \hat{g}_t \odot \hat{g}_t \\ \hat{s}_t &\leftarrow \frac{s_t}{1 - \rho_1^t} \\ \hat{r}_t &\leftarrow \frac{r_t}{1 - \rho_2^t} \\ \Delta \theta &\leftarrow - \frac{\eta_0}{\delta + \sqrt{\hat{r}_t}} \odot \hat{s}_t \\ \theta_{t+1} &\leftarrow \theta_t + \Delta \theta\end{aligned}$$



# Algorithm Comparison on MNIST Dataset



# Algorithm Comparison on MNIST Dataset



# Algorithm Comparison on MNIST Dataset

Training Accuracy vs. Number of examples seen

